

## 02\_linear\_regression\_grad\_desc

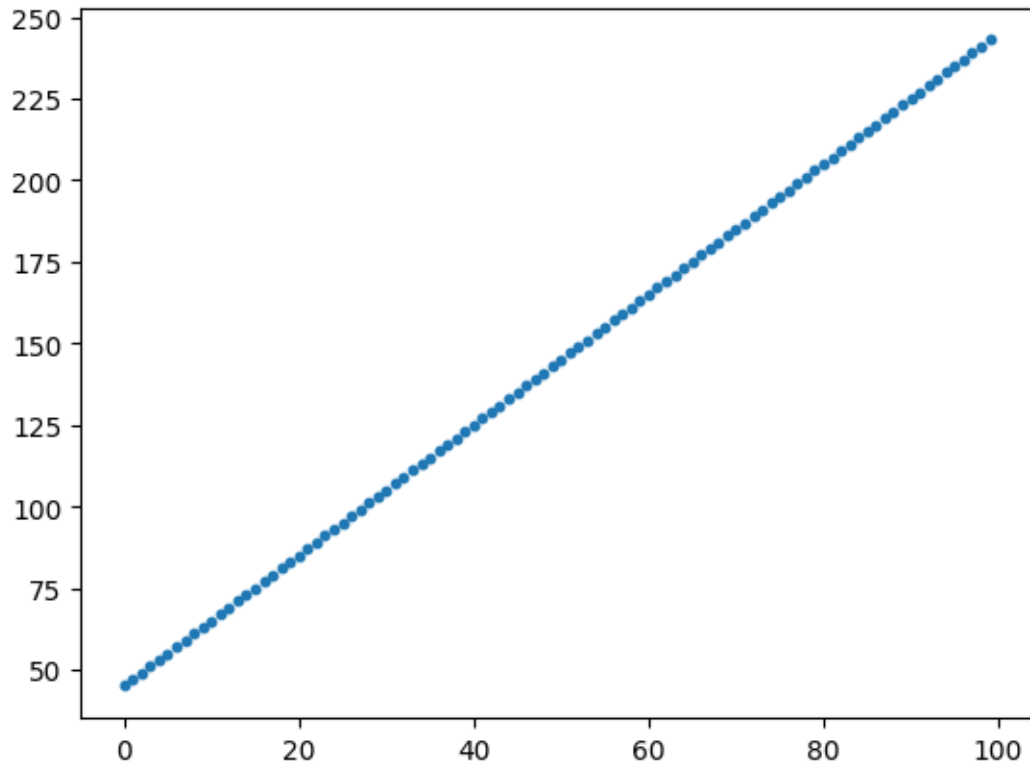
March 10, 2025

```
[1]: import numpy as np  
import matplotlib.pyplot as plt
```

```
[2]: def linear_function(a, x, b):  
    return a*x + b
```

```
[3]: a = 2  
x = np.arange(100)  
b = 45  
y = linear_function(a, x, b)
```


```
[4]: plt.scatter(x, y, s=10)  
plt.show()
```



```
[5]: # https://medium.com/@ms\_somanna/
      ↪ guide-to-adding-noise-to-your-data-using-python-and-numpy-c8be815df524

      noise = np.random.normal(0, 10, len(x))
      y_noisy = y + noise

[6]: plt.scatter(x, y_noisy, s=10)
      plt.show()
```



## 1 Gradient Descent Derivation

The goal of gradient descent is to minimize the Mean Squared Error (MSE) loss function for a linear model:

$$y = \beta_1 x + \beta_0$$

### 1.1 Loss Function

The MSE loss function is defined as:

$$J(\beta_0, \beta_1) = \frac{1}{n} \sum_{i=1}^n (y_i - (\beta_1 x_i + \beta_0))^2$$

where: -  $n$  is the number of data points, -  $y_i$  is the actual observed value, -  $x_i$  is the feature (independent variable), -  $\beta_0$  is the intercept, -  $\beta_1$  is the slope.

## 1.2 Gradient Computation

To update the parameters  $\beta_0$  and  $\beta_1$ , we compute the partial derivatives of  $J(\beta_0, \beta_1)$  with respect to each parameter.

### 1.2.1 Partial Derivative with Respect to $\beta_1$

$$\frac{\partial J}{\partial \beta_1} = \frac{2}{n} \sum_{i=1}^n (y_i - (\beta_1 x_i + \beta_0))(-x_i) = -\frac{2}{n} \sum_{i=1}^n x_i (y_i - (\beta_1 x_i + \beta_0))$$

### 1.2.2 Partial Derivative with Respect to $\beta_0$

$$\frac{\partial J}{\partial \beta_0} = \frac{2}{n} \sum_{i=1}^n (y_i - (\beta_1 x_i + \beta_0))(-1) = -\frac{2}{n} \sum_{i=1}^n (y_i - (\beta_1 x_i + \beta_0))$$

## 1.3 Gradient Descent Updates

Using gradient descent, we update the parameters iteratively as follows:

$$\beta_1 := \beta_1 - \alpha \frac{\partial J}{\partial \beta_1}$$

$$\beta_0 := \beta_0 - \alpha \frac{\partial J}{\partial \beta_0}$$

where  $\alpha$  is the learning rate, controlling the step size in the descent process.

```
[7]: def gradient_descent(x, y, learning_rate=0.0001, epochs=100000):
    n = len(x)
    beta_0, beta_1 = 0, 0 # Initialize parameters

    for _ in range(epochs):
        y_pred = beta_1 * x + beta_0
        error = y_pred - y

        grad_beta_1 = (2/n) * np.sum(error * x)
        grad_beta_0 = (2/n) * np.sum(error)

        beta_1 -= learning_rate * grad_beta_1
        beta_0 -= learning_rate * grad_beta_0

    return beta_1, beta_0
```

```
# Run gradient descent  
beta_1, beta_0 = gradient_descent(x, y_noisy)  
print(f"Estimated beta_1: {beta_1}, Estimated beta_0: {beta_0}")
```

Estimated beta\_1: 1.9844676182020553, Estimated beta\_0: 46.45394715662439

```
[8]: plt.scatter(x, y_noisy, s=10)  
plt.plot(beta_0 + beta_1 * x)  
plt.show()
```

